

Retrieval-Augmented Generation

Retrieval-Augmented Generation (RAG) combines a retrieval system with a generative language model. Instead of relying solely on the model's internal knowledge, RAG retrieves relevant documents from an external knowledge base and uses them as context. This approach reduces hallucination and allows models to answer questions using up-to-date or domain-specific information not seen during training.